

Working with Colors

```
> library(swirl)
```

```
| Hi! I see that you have some variables saved in your workspace. To keep things running smoothly, I  
| recommend you clean up before starting swirl.
```

```
| Type ls() to see a list of the variables in your workspace. Then, type rm(list=ls()) to clear your  
| workspace.
```

```
| Type swirl() when you are ready to begin.
```

```
> swirl()
```

```
| Welcome to swirl! Please sign in. If you've been here before, use the same name as you did then. If  
| you are new, call yourself something unique.
```

```
What shall I call you? Stephen
```

```
| Would you like to continue with one of these lessons?
```

```
1: R Programming Workspace and Files  
2: No. Let me start something new.
```

```
Selection: 2
```

```
| Please choose a course, or type 0 to exit swirl.
```

```
1: Data Analysis  
2: Exploratory Data Analysis  
3: Getting and Cleaning Data  
4: Mathematical Biostatistics Boot Camp  
5: Open Intro  
6: R Programming  
7: Regression Models  
8: Statistical Inference  
9: Take me to the swirl course repository!
```

```
Selection: 2
```

```
| Please choose a lesson, or type 0 to return to course menu.
```

```
1: Principles of Analytic Graphs    2: Exploratory Graphs  
3: Graphics Devices in R          4: Plotting Systems  
5: Base Plotting System           6: Lattice Plotting System  
7: Working with Colors            8: GGPlot2 Part1  
9: GGPlot2 Part2                 10: GGPlot2 Extras  
11: Hierarchical Clustering       12: K Means Clustering  
13: Dimension Reduction          14: Clustering Example  
15: CaseStudy
```

```
Selection: 7
```

| Attempting to load lesson dependencies...

| Package 'jpeg' loaded correctly!

| Package 'RColorBrewer' loaded correctly!

| Package 'datasets' loaded correctly!

|
| 0%

| working_with_Colors. (Slides for this and other Data Science courses may be found at github

| <https://github.com/DataScienceSpecialization/courses/>. If you care to use them, they must be

| downloaded as a zip file and viewed locally. This lesson corresponds to
| 04_ExploratoryAnalysis/Colors.)

...

|=
| 1%

| This lesson is about using colors in R. It really supplements the lessons on plotting with the base

| and lattice packages which contain functions that are able to take the argument col. We'll discuss

| ways to set this argument more colorfully.

...

|===
| 3%

| Of course, color choice is secondary to your data and how you analyze it, but effectively using

| colors can enhance your plots and presentations, emphasizing the important points you're trying to

| convey.

...

|====
| 4%

| The motivation for this lesson is that the default color schemes for most plots in R are not

| optimal. Fortunately there have been recent developments to improve the handling and specification

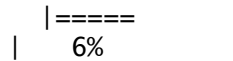
| of colors in plots and graphs. We'll cover some functions in R as well as in external packages that

| are very handy. If you know how to use some of these then you'll have more options when you create

| your displays.

...

|=====

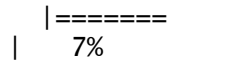


| 6%

| We'll begin with a motivating example - a typical R plot using 3 default colors.

...

|=====



| 7%

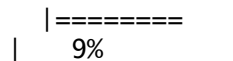
| According to the plot, what is color 2?

1: Red
2: Blue
3: Empty black circles
4: Green

Selection: 1

| Excellent job!

|=====

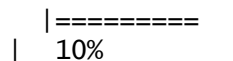


| 9%

| So these are the first 3 default values. If you were plotting and just specified `col=c(1:3)` as one of your arguments, these are colors you'd get. Maybe you like them, but they might not be the best choice for your application.

...

|=====

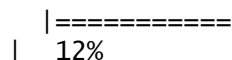


| 10%

| To show you some options, here's a display of two color palettes that come with the `grDevices` package available to you. The left shows you some colors from the function `heat.colors`. Here low values are represented in red and as the values increase the colors move through yellow towards white. This is consistent with the physical properties of fire. The right display is from the function `topo.colors` which uses topographical colors ranging from blue (low values) towards brown (higher values).

...

|=====



| 12%

| So we'll first discuss some functions that the `grDevices` package offers. The function `colors()` lists the names of 657 predefined colors you can use in any plotting function. These names are

| returned as strings. Run the R command `sample` with `colors()` as its first argument and 10 as its second to give you an idea of the choices you have.

```
> sample(colors(),10)
[1] "lightyellow2" "ivory2" "palevioletred1" "grey100" "grey76"
[6] "coral" "mediumpurple2" "rosybrown4" "grey94" "chartreuse"
```

| That's the answer I was looking for.

```
|=====
| 13%
```

| We see a lot of variety in the colors, some of which are names followed by numbers indicating that there are multiple forms of that particular color.

...

```
|=====
| 14%
```

| So you're free to use any of these 600+ colors listed by the `colors` function. However, two additional functions from `grDevices`, `colorRamp` and `colorRampPalette`, give you more options. Both of these take color names as arguments and use them as "palettes", that is, these argument colors are blended in different proportions to form new colors.

...

```
|=====
| 16%
```

| The first, `colorRamp`, takes a palette of colors (the arguments) and returns a function that takes values between 0 and 1 as arguments. The 0 and 1 correspond to the extremes of the color palette. Arguments between 0 and 1 return blends of these extremes.

...

```
|=====
| 17%
```

| Let's see what this means. Assign to the variable `pal` the output of a call to `colorRamp` with the single argument, `c("red","blue")`.

```
> pal <- colorRamp(c("red","blue"))
```

| Excellent job!

```
|=====
| 19%
```

| We don't see any output, but R has created the function pal which we can call with a single argument between 0 and 1. Call pal now with the argument 0.

```
> View(pal)
```

| Almost! Try again. Or, type info() for more options.

| Type pal(0) at the command prompt.

```
> pal(0)
      [,1] [,2] [,3]
[1,] 255   0   0
```

| Excellent work!

```
|=====
| 20%
```

| We see a 1 by 3 array with 255 as the first entry and 0 in the other 2. This 3 long vector corresponds to red, green, blue (RGB) color encoding commonly used in television and monitors. In R, 24 bits are used to represent colors. Think of these 24 bits as 3 sets of 8 bits, each of which represents an intensity for one of the colors red, green, and blue.

...

```
|=====
| 22%
```

| The 255 returned from the pal(0) call corresponds to the largest possible number represented with 8 bits, so the vector (255,0,0) contains only red (no green or blue), and moreover, it's the highest possible value of red.

...

```
|=====
| 23%
```

| Given that you created pal with the palette containing "red" and "blue", what color do you think will be represented by the vector that pal(1) returns? Recall that pal will only take arguments between 0 and 1, so 1 is the largest argument you can pass it.

```
1: green
2: blue
3: red
4: yellow
```

Selection: 1

| One more time. You can do it!

| If the one extreme (0) returned red, it makes sense that pal(1) will return a vector representing
| blue.

1: blue
2: yellow
3: green
4: red

Selection: 1

| You are amazing!

|=====

| 25%

| Check your answer now by calling pal with the argument 1.

```
> pal(1)
      [,1] [,2] [,3]
[1,]    0    0 255
```

| Keep working like that and you'll get there!

|=====

| 26%

| You see the vector (0,0,255) which represents the highest intensity of blue
|. What vector do you
| think the call pal(.5) will return?

1: (255,255,255)
2: (127.5,0,127.5)
3: (0,255,0)
4: (255,0,255)

Selection: 3

| That's not exactly what I'm looking for. Try again.

| The correct answer should be halfway between (255,0,0) and (0,0,255). which
| is the only choice that
| is the average (mean) of these two?

1: (127.5,0,127.5)
2: (0,255,0)
3: (255,0,255)
4: (255,255,255)

Selection: 1

| Your dedication is inspiring!

```
|=====
| 28%
```

| The function `pal` can take more than one argument. It returns one 3-long (or 4-long, but more about this later) vector for each argument. To see this in action, call `pal` with the argument `seq(0,1,len=6)`.

```
> pal(seq(0,1,len=6))
Error in seq.default(0, 1, len = 6) : object 'len' not found
> pal(seq(0,1,len=6))
      [,1] [,2] [,3]
[1,]  255    0    0
[2,]  204    0   51
[3,]  153    0  102
[4,]  102    0  153
[5,]   51    0  204
[6,]    0    0  255
```

| That's a job well done!

```
|=====
| 29%
```

| Six vectors (each of length 3) are returned. The i -th vector is identical to output that would be returned by the call `pal(i/5)` for $i=0,\dots,5$. We see that the i -th row (for $i=1,\dots,6$) differs from the $(i-1)$ -st row in the following way. Its red entry is $51 = 255/5$ points lower and its blue entry is 51 points higher.

...

```
|=====
| 30%
```

| So `pal` creates colors using the palette we specified when we called `colorRamp`. In this example none of `pal`'s outputs will ever contain green since it wasn't in our initial palette.

...

```
|=====
| 32%
```

| We'll turn now to `colorRampPalette`, a function similar to `colorRamp`. It also takes a palette of colors and returns a function. This function, however, takes integer arguments (instead of numbers between 0 and 1) and returns a vector of colors each of which is a blend of colors of the original palette.

...

```
|=====
| 33%
```

| The argument you pass to the returned function specifies the number of colors you want returned.
| Each element of the returned vector is a 24 bit number, represented as 6 hexadecimal characters,
| which range from 0 to F. This set of 6 hex characters represents the intensities of red, green, and
| blue, 2 characters for each color.

...

```
|=====
| 35%
```

| To see this better, assign to the variable p1 the output of a call to colorRampPalette with the
| single argument, c("red","blue"). We'll compare it to our experiments using colorRamp.

```
> p1 <- colorRampPalette(c("red","blue"))
```

| Keep up the great work!

```
|=====
| 36%
```

| Now call p1 with the argument 2.

```
> p1(2)
[1] "#FF0000" "#0000FF"
```

| Nice work!

```
|=====
| 38%
```

| We see a 2-long vector is returned. The first entry FF0000 represents red. The FF is hexadecimal
| for 255, the same value returned by our call pal(0). The second entry 0000FF represents blue, also
| with intensity 255.

...

```
|=====
| 39%
```

| Now call p1 with the argument 6. Let's see if we get the same result as we did when we called pal
| with the argument seq(0,1,len=6).

```
> p1(6)
[1] "#FF0000" "#CC0033" "#990066" "#650099" "#3200CC" "#0000FF"
```


| You nailed it! Good job!

|=====

| 41%

| Now we get the 6-long vector (FF0000, CC0033, 990066, 650099, 3200CC, 0000FF). We see the two ends (FF0000 and 0000FF) are consistent with the colors red and blue. How about CC0033? Type 0xcc or 0xCC at the command line to see the decimal equivalent of this hex number. You must include the 0 before the x to specify that you're entering a hexadecimal number.

```
> 0xcc
[1] 204
```

| You're the best!

|=====

| 42%

| So 0xCC equals 204 and we can easily convert hex 33 to decimal, as in $0x33 = 3 \times 16 + 3 = 51$. These were exactly the numbers we got in the second row returned from our call to `pal(seq(0,1,len=6))`. We see that 4 of the 6 numbers agree with our earlier call to `pal`. Two of the 6 differ slightly.

...

|=====

| 43%

| We can also form palettes using colors other than red, green and blue. Form a palette, `p2`, by calling `colorRampPalette` with the colors "red" and "yellow". Remember to concatenate them into a single argument.

```
> p2 <- colorRampPalette(c("red","yellow"))
```

| You got it right!

|=====

| 45%

| Now call `p2` with the argument 2. This will show us the two extremes of the blends of colors we'll get.

```
> p2(2)
[1] "#FF0000" "#FFFF00"
```

| All that hard work is paying off!

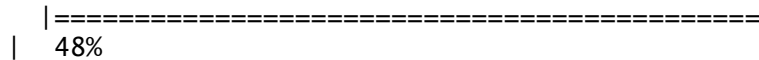
|=====

| 46%

| Not surprisingly the first color we see is FF0000, which we know represents red. The second color
| returned, FFFF00, must represent yellow, a combination of full intensity red and full intensity
| green. This makes sense, since yellow falls between red and green on the color wheel as we see
| here. (We borrowed this image from lucaskrech.com.)

...

|=====



| 48%

| Let's now call p2 with the argument 10. This will show us how the two extremes, red and yellow, are
| blended together.

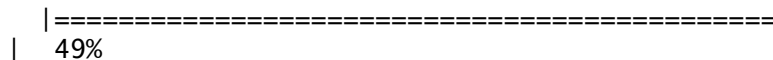
>

> p2(10)

```
[1] "#FF0000" "#FF1C00" "#FF3800" "#FF5500" "#FF7100" "#FF8D00" "#FFAA00" "#FFC600" "#FFE200"
[10] "#FFFF00"
```

| Excellent job!

|=====

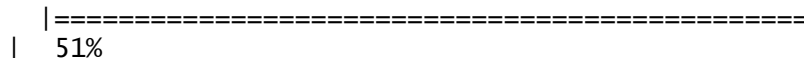


| 49%

| So we see the 10-long vector. For each element, the red component is fixed at FF, and the green
| component grows from 00 (at the first element) to FF (at the last).

...

|=====



| 51%

| This is all fine and dandy but you're probably wondering when you can see how all these colors show
| up in a display. We copied some code from the R documentation pages (color.scale if you're
| interested) and created a function for you, showMe. This takes as an argument, a color vector,
| which as you know, is precisely what calls to p1 and p2 return to you. Call showMe now with p1(20).

> View(showMe)

| Nice try, but that's not exactly what I was hoping for. Try again. Or, type info() for more
| options.

| Type showMe(p1(20)) at the command prompt.

> View(showMe)

| Not quite! Try again. Or, type info() for more options.

| Type showMe(p1(20)) at the command prompt.

```
> showMe(p1(20))
```

| That's a job well done!

```
|=====
| 52%
```

| We see the interpolated palette here. Low values in the lower left corner are red and as you move to the upper right, the colors move toward blue. Now call showMe with p2(20) as its argument.

```
> showMe(p2(20))
```

| You are doing so well!

```
|=====
| 54%
```

| Here we see a similar display, the colors moving from red to yellow, the base colors of our p2 palette. For fun, see what p2(2) looks like using showMe.

```
> p2(2)
```

```
[1] "#FF0000" "#FFFF00"
```

| Not quite, but you're learning! Try again. Or, type info() for more options.

| Type showMe(p2(2))) at the command prompt.

```
> showMe((p2(2)))
```

| Keep trying! Or, type info() for more options.

| Type showMe(p2(2))) at the command prompt.

```
> showMe(p2(2))
```

| You're the best!

```
|=====
| 55%
```

| A much more basic pattern, simple but elegant.

...

```
|=====
| 57%
```

| We mentioned before that colorRamp (and colorRampPalette) could return a 3 or 4 long vector of

| colors. We saw 3-long vectors returned indicating red, green, and blue intensities. What would the 4th entry be?

...

|=====

| 58%

| We'll answer this indirectly. First, look at the function p1 that colorRampPalette returned to you.
| Just type p1 at the command prompt.

```
> p1
function (n)
{
  x <- ramp(seq.int(0, 1, length.out = n))
  if (ncol(x) == 4L)
    rgb(x[, 1L], x[, 2L], x[, 3L], x[, 4L], maxColorValue = 255)
  else rgb(x[, 1L], x[, 2L], x[, 3L], maxColorValue = 255)
}
<bytecode: 0x000000000579dbb8>
<environment: 0x0000000005889240>
```

| You're the best!

|=====

| 59%

| We see that p1 is a short function with one argument, n. The argument n is used as the length in a call to the function seq.int, itself an argument to the function ramp. We can infer that ramp is just going to divide the interval from 0 to 1 into n pieces.

...

|=====

| 61%

| The heart of p1 is really the call to the function rgb with either 4 or 5 arguments. Use the ?fun construct to look at the R documentation for rgb now.

```
> ?fun
No documentation for 'fun' in specified packages and libraries:
you could try '??fun'
```

| Give it another try. Or, type info() for more options.

| Type ?rgb at the command prompt.

```
> ?rgb
```

| Keep up the great work!

```
|=====
| 62%
```

```
| We see that rgb is a color specification function that can be used to produce any color with red,
| green, blue proportions. We see the maxColorValue is 1 by default, so if we called rgb with values
| for red, green and blue, we would specify numbers at most 1 (assuming we didn't change the default
| for maxColorValue). According to the documentation, what is the maximum number of arguments rgb can
| have?
```

```
1: 6
2: 3
3: 5
4: 4
```

selection: 1

```
| Keep working like that and you'll get there!
```

```
|=====
| 64%
```

```
| So the fourth argument is alpha which can be a logical, i.e., either TRUE or FALSE, or a numerical
| value. Create the function p3 now by calling colorRampPalette with the colors blue and green
| (remember to concatenate them into a single argument) and setting the alpha argument to .5.
```

```
> p3 <- colorRampPalette(c("blue","green"),alpha=.5)
```

```
| All that hard work is paying off!
```

```
|=====
| 65%
```

```
| Now call p3 with the argument 5.
```

```
> p3(5)
[1] "#0000FFFF" "#003FBFFF" "#007F7FFF" "#00BF3FFF" "#00FF00FF"
```

```
| You are doing so well!
```

```
|=====
| 67%
```

```
| We see that in the 5-long vector that the call returned, each element has 32 bits, 4 groups of 8
| bits each. The last 8 bits represent the value of alpha. Since it was NOT ZERO in the call to
| colorRampPalette, it gets the maximum FF value. (The same result would happen if alpha had been set
| to TRUE.) When it was 0 or FALSE (as in previous calls to colorRampPalette) it was given the value
```

| 00 and wasn't shown. The leftmost 24 bits of each element are the same RGB encoding we previously saw.

...

|=====

| 68%

| So what is alpha? Alpha represents an opacity level, that is, how transparent should the colors be.

| We can add color transparency with the alpha parameter to calls to rgb. We haven't seen any

| examples of this yet, but we will now.

...

|=====

| 70%

| We generated 1000 random normal pairs for you in the variables x and y. We'll plot them in a

| scatterplot by calling plot with 4 arguments. The variables x and y are the first 2. The third is

| the print character argument pch. Set this equal to 19 (filled circles). The final argument is col

| which should be set equal to a call to rgb. Give rgb 3 arguments, 0, .5, and .5.

```
> plot(x,y,pch=19,col=rgb(0,.5,.5))
```

| Keep up the great work!

|=====

| 71%

| Well this picture is okay for a scatterplot, a nice mix of blue and green, but it really doesn't

| tell us too much information in the center portion, since the points are so thick there. We see

| there are a lot of points, but is one area more filled than another? We can't really discriminate

| between different point densities. This is where the alpha argument can help us. Recall your plot

| command (use the up arrow) and add a 4th argument, .3, to the call to rgb. This will be our value

| for alpha.

```
> plot(x,y,pch=19,col=rgb(0,.5,.5,alpha = 0.3))
```

| You are really on a roll!

|=====

| 72%

| Clearly this is better. It shows us where, specifically, the densest areas of the scatterplot

| really are.

...

|=====

| 74%

| Our last topic for this lesson is the RColorBrewer Package, available on CRAN, that contains
| interesting and useful color palettes, of which there are 3 types, sequential, divergent, and
| qualitative. Which one you would choose to use depends on your data.

...

|=====

| 75%

| Here's a picture of the palettes available from this package. The top section shows the sequential
| palettes in which the colors are ordered from light to dark. The divergent palettes are at the
| bottom. Here the neutral color (white) is in the center, and as you move from the middle to the two
| ends of each palette, the colors increase in intensity. The middle display shows the qualitative
| palettes which look like collections of random colors. These might be used to distinguish factors
| in your data.

...

|=====

| 77%

| These colorBrewer palettes can be used in conjunction with the colorRamp() and colorRampPalette()
| functions. You would use colors from a colorBrewer palette as your base palette, i.e., as arguments
| to colorRamp or colorRampPalette which would interpolate them to create new colors.

...

|=====

| 78%

| As an example of this, create a new object, cols by calling the function brewer.pal with 2
| arguments, 3 and "BuGn". The string "BuGn" is the second last palette in the sequential display.
| The 3 tells the function how many different colors we want.

```
> cols <- brewer.pal(3, "BuGn")
```

| That's the answer I was looking for.

```

|=====
| 80%

| Use showMe to look at cols now.

> showMe(cols)

| Your dedication is inspiring!

= |=====
  | 81%

| We see 3 colors, mixes of blue and green. Now create the variable pal by ca
|lling colorRampPalette
| with cols as its argument.

> pal <- colorRampPalette(cols)

| That's the answer I was looking for.

== |=====
   | 83%

| The call showMe(pal(3)) would be identical to the showMe(cols) call. So use
| showMe to look at
| pal(20).

> showMe(pal(20))

| You are quite good my friend!

=== |=====
    | 84%

| Now we can use the colors in pal(20) to display topographic information on
| Auckland's Maunga Whau
| Volcano. R provides this information in a matrix called volcano which is in
| cluded in the package
| datasets. Call the R function image with volcano as its first argument and
| col set equal to
| pal(20) as its second.

> image(volcano,col = pal(20))

| You got it right!

===== |=====
      | 86%

| We see that the colors here of the sequential palette clue us in on the top
| ography. The darker
| colors are more concentrated than the lighter ones. Just for fun, recall yo
| ur last command calling
| image and instead of pal(20), use p1(20) as the second argument.

> image(volcano,col = p1(20))

```


| All that hard work is paying off!

=====|=====87%

| Not as nice a picture since the palette isn't as well suited to this data,
but that's okay. It's
| review time!!!!

...

=====|=====88%

| True or False? Careful use of colors in plots/maps/etc. can make it easier
for the reader to
| understand what points you're trying to convey.

- 1: True
- 2: False

selection: 1

| Keep up the great work!

=====|=====90%

| Which of the following is an R package that provides color palettes for sequential, categorical,
| and diverging data?

- 1: RColorVintner
- 2: RColorBrewer
- 3: RColorStewer
- 4: RColorBluer

selection: 2

| That's correct!

=====|=====91%

| True or False? The colorRamp and colorRampPalette functions can be used in
conjunction with color
| palettes to connect data to colors.

- 1: False
- 2: True

selection: 2

| You nailed it! Good job!

=====|=====93%

| True or False? Transparency can NEVER be used to clarify plots with many points

1: False
2: True

selection: 1

| Perseverance, that's the answer.

|=====

===== | 94%

| True or False? The call p7 <- colorRamp("red","blue") would work (i.e., not generate an error).

1: True
2: False

selection: 2

| All that hard work is paying off!

|=====

===== | 96%

| True or False? The function colors returns only 10 colors.

1: True
2: False

selection: 2

| You're the best!

|=====

===== | 97%

| Transparency is determined by which parameter of the rgb function?

1: alpha
2: beta
3: gamma
4: delta
5: it's all Greek to me

selection: 1

| You got it!

|=====

===== | 99%

| Congratulations! we hope this lesson didn't make you see red. We're green with envy that you blue through it.

...


|=====| 100%

| would you like to receive credit for completing this course on Coursera.org ?



1: Yes


2: No

p1(20)




$p_2(20)$





cols



pal(20)

