

Clustering Example

```
> library(swirl)
```

```
| Hi! Type swirl() when you are ready to begin.
```

```
> swirl()
```

```
| Welcome to swirl! Please sign in. If you've been here before, use the same name as you did then  
are new,  
| call yourself something unique.
```

```
What shall I call you? Stephen
```

```
| Please choose a course, or type 0 to exit swirl.
```

```
1: Data Analysis  
2: Exploratory Data Analysis  
3: Getting and Cleaning Data  
4: Mathematical Biostatistics Boot Camp  
5: Open Intro  
6: R Programming  
7: Regression Models  
8: Statistical Inference  
9: Take me to the swirl course repository!
```

```
Selection: 2
```

```
| Please choose a lesson, or type 0 to return to course menu.
```

```
1: Principles of Analytic Graphs    2: Exploratory Graphs          3: Graphics Devices in R  
4: Plotting Systems              5: Base Plotting System       6: Lattice Plotting System  
7: Working with Colors           8: GGPlot2 Part1             9: GGPlot2 Part2  
10: GGPlot2 Extras              11: Hierarchical Clustering  12: K Means Clustering  
13: Dimension Reduction         14: Clustering Example       15: CaseStudy
```

```
Selection: 14
```

```
| Attempting to load lesson dependencies...
```

```
| Package 'fields' loaded correctly!
```

```
| Package 'jpeg' loaded correctly!
```

```
| Package 'datasets' loaded correctly!
```

```
|  
| 0%
```

```
| Clustering_Example. (Slides for this and other Data Science courses may be found at github  
| https://github.com/DataScienceSpecialization/courses/. If you care to use them, they must be down-  
as a zip  
| file and viewed locally. This lesson corresponds to 04_ExploratoryAnalysis/clusteringExample.)
```

```
...
```

```
|  
| ==  
| 2%
```

```
| In this lesson we'll apply some of the analytic techniques we learned in this course to data from  
university  
| of California, Irvine. Specifically, the data we'll use is from UCI's Center for Machine Learning  
Intelligent  
| Systems. You can find out more about the data at  
| http://archive.ics.uci.edu/ml/datasets/Human+Activity+Recognition+Using+Smartphones. As this ad-  
dicates,  
| the data involves smartphones and recognizing human activity. Cool, right?
```

```
...
```

```
|===  
| 3%
```

```
| Our goal is to show you how to use exploratory data analysis to point you in fruitful directions  
earch, that  
| is, towards answerable questions. Exploratory data analysis is a "rough cut" or filter which helps  
to find the  
| most beneficial areas of questioning so you can set your priorities accordingly.
```

```
...
```

```
|=====  
| 5%
```

```
| We also hope to show you that "real-world" research isn't always neat and well-defined like textbook  
estions  
| with clearcut answers.
```

```
...
```

```
|=====  
| 6%
```

```
| We've loaded data from this study for you in a matrix called ssd. Run the R command dim now to see the  
| dimensions.
```

```
> dim(ssd)  
[1] 7352 563
```

```
| Keep up the great work!
```

```
|=====  
| 8%
```

```
| Wow - ssd is pretty big, 7352 observations, each of 563 variables. Don't worry we'll only use a small  
ortion of  
| this "Human Activity Recognition database".
```

```
...
```

```
|=====  
| 9%
```

```
| The study creating this database involved 30 volunteers "performing activities of daily living" while  
| carrying a waist-mounted smartphone with embedded inertial sensors. ... Each person performed several  
ities ...  
| wearing a smartphone (Samsung Galaxy S II) on the waist. ... The experiments have been video-recorded  
o label the  
| data manually. The obtained dataset has been randomly partitioned into two sets, where 70% of the  
nteers was  
| selected for generating the training data and 30% the test data."
```

```
...
```

```
|=====  
| 11%
```

```
| Use the R command names with just the last two columns (562 and 563) of ssd to see what data the  
in.
```

```
> names(ssd[562:563])  
[1] "subject" "activity"
```

```
| You got it right!
```

```
|=====  
| 12%
```

| These last 2 columns contain subject and activity information. We saw above that the gathered data has
"been
| randomly partitioned into two sets, where 70% of the volunteers was selected for generating the training
g data and
| 30% the test data." Run the R command table with ssd\$subject as its argument to see if the data set
contains
| training or test data.

```
>  
> table(ssd$subject)
```

```
 1    3    5    6    7    8   11   14   15   16   17   19   21   22   23   25   26   27   28   29   30  
347 341 302 325 308 281 316 323 328 366 368 360 408 321 372 409 392 376 382 344 383
```

| Your dedication is inspiring!

```
|=====
| 14%
```

| From the number of subjects, would you infer that ssd contains training or test data?

```
1: test
2: training
```

Selection: 2

| You got it right!

```
|=====
| 16%
```

| So ssd contains only training data. If you ran the R command sum with table(ssd\$subject) as its argument, what
t, what
| would the number you get back represent?

```
1: the number of rows in ssd
2: Huh?
3: the number of columns in ssd
4: the number of rows and columns of ssd
```

Selection: 1

| That's correct!

```
|=====
| 17%
```

| Try it now (running sum on table(ssd\$subject)) to see if you get 7352, the number of rows in ssd. What's the
result.

```
> sum(table(ssd$subject))
[1] 7352
```

| You are quite good my friend!

```
|=====
| 19%
```

| So we're looking at training data from a machine learning repository. We can infer that this data is for
pposed to
| train machines to recognize activity collected from the accelerometers and gyroscopes built into
artphones
| that the subjects had strapped to their waists. Run the R command table on ssd\$activity to see what activities
ivities
| have been characterized by this data.

```
> table(ssd$activity)
```

```
laying  sitting  standing    walk  walkdown  walkup
 1407    1286    1374    1226      986    1073
```

| You nailed it! Good job!

|=====

| 20%

| We have 6 activities, 3 passive (laying, standing and sitting) and 3 active which involve walking or running.

| You ran the R command `sum(table(ssd$activity))` as its argument, what would the number you get back represent?

- 1: Huh?
- 2: the number of rows in `ssd`
- 3: the number of columns in `ssd`
- 4: the number of rows and columns of `ssd`

Selection: 2

| You nailed it! Good job!

|=====

| 22%

| Because it's training data, each row is labeled with the correct activity (from the 6 possibilities).

| We're interested in questions like "Is the correlation between the measurements and activities good enough to train a machine?" so that we can use a set of

| 561 measurements, would a trained machine be able to determine which of the 6 activities the person was doing?"

...

|=====

| 23%

| First, let's massage the data a little so it's easier to work with. We've already run the R command `ssd$activity` on

| the data so that activities are factors. This will let us color code them when we generate plots. Let's look at

| only the first subject (numbered 1). Create the variable `sub1` by assigning to it the output of the R command

| `subset(ssd, subject == 1)`, as the second.

> `sub1 <- subset(ssd, subject == 1)`

| You are doing so well!

|=====

| 25%

| Look at the dimensions of `sub1` now.

> `dim(sub1)`

[1] 347 563

| All that hard work is paying off!

|=====

| 27%

| So `sub1` has fewer than 400 rows now, but still a lot of columns which contain measurements. Use the first

| 12 columns of `sub1` to see what kind of data we have.

> `names(sub1[1:12])`

[1] "tBodyAcc.mean...X" "tBodyAcc.mean...Y" "tBodyAcc.mean...Z" "tBodyAcc.std...X" "tBodyAcc.std...Y" "tBodyAcc.std...Z" "tBodyAcc.mad...X" "tBodyAcc.mad...Y" "tBodyAcc.mad...Z" "tBodyAcc.max...X" "tBodyAcc.max...Y" "tBodyAcc.max...Z"

| You're the best!

```

|=====
| 28%

| We see X, Y, and Z (3 dimensions) of different aspects of body acceleration measurements, such as
and
| standard deviation. Let's do some comparisons of activities now by looking at plots of mean body
acceleration in
| the X and Y directions. Call the function myedit with the string "showXY.R" to see the code generating
the plots.
| Make sure your cursor is back in the console window before you hit any more buttons.

> myedit("showXY.R")

| You're the best!

|=====
| 30%

| You see both the code and its output! The plots are a little squished, but we see that the activities
related to walking (shown in the two blues and magenta) show more variability than the passive activities
(shown
in black, red, and green), particularly in the X dimension.

...

|=====
| 31%

| The colors are a little hard to distinguish. Just for fun, call the function showMe (we used it in the
Working_with_Colors lesson) which displays color vectors. Use the vector 1:6 as its argument and hopefully
this will clarify the colors you see in the XY comparison plot.

> showMe(1:6)

| You are doing so well!

|=====
| 33%

| Nice! We just wanted to show you the beauty and difference in colors. The colors at the bottom, black,
red and green, mark the passive activities, while the true blues and magenta near the top show the walking
activities. Let's try clustering to see if we can distinguish the activities more.

...

|=====
| 34%

| We'll still focus on the 3 dimensions of mean acceleration. (The plot we just saw looked at the first 2
dimensions.) Create a distance matrix, mdist, of the first 3 columns of sub1, by using the R command dist.
Use the x[,1:3] notation to specify the columns.

> mdist <- dist(sub1[,1:3])

| You got it!

|=====
| 36%

| Now create the variable hclustering by calling the R command hclust and passing it mdist as an argument.
This will use the Euclidean distance as its default metric.

> hclustering <- hclust(mdist)

| All that hard work is paying off!

|=====
| 38%

| Now call the pretty plotting function (which we've already sourced) myplotclust with 2 arguments.

```

```

| first is hclustering, and the second is the argument lab.col set equal to unclass(sub1$activity)
> myplclust(hclustering, lab.col = unclass(sub1$activity))

| All that practice is paying off!

|=====
| 39%

| Well that dendrogram doesn't look too helpful, does it? There's no clear grouping of colors, except
| that active colors (blues and magenta) are near each other as are the passive (black, red, and
| green). So average acceleration doesn't tell us much. How about maximum acceleration? Let's look
| that for the first subject (in our array sub1) for the X and Y dimensions. These are in column 10
| and 11.

...

|=====
| 41%

| Here they are plotted side by side, X dimension on the left and Y on the right. The x-axis of each
| shows the 300+ observations and the y-axis indicates the maximum acceleration.

...

|=====
| 42%

| From the 2 plots, what separation, if any, do you see?

1: passive activities generate the most acceleration
2: there is no pattern
3: laying generates the most acceleration in the X dimension
4: passive activities mostly fall below the walking activities

Selection: 4

| Perseverance, that's the answer.

|=====
| 44%

| Finally we're seeing something vaguely interesting! Let's focus then on the 3 dimensions of maximum
| acceleration, stored in columns 10 through 12 of sub1. Create a new distance matrix, mdist, of
| these 3 columns of sub1, by using the R command dist. Again, use the x[,10:12] notation to catch
| the columns.

> mdist <- dist(sub1[,10:12])

| You nailed it! Good job!

|=====
| 45%

| Now create the variable hclustering by calling hclust with mdist as the argument.

> hclustering <- hclust(mdist)

| You're the best!

|=====
| 47%

| Again, call the myplclust with 2 arguments. The first is hclustering, and the second is the
| argument lab.col set equal to unclass(sub1$activity).

> myplclust(hclustering, lab.col = unclass(sub1$activity))

| Great job!

```

```
|=====
| 48%
```

```
| Now we see clearly that the data splits into 2 clusters, active and passive activities. Moreover,
| the light blue (walking down) is clearly distinct from the other walking activities. The dark blue
| (walking level) also seems to be somewhat clustered. The passive activities, however, seem all
| jumbled together with no clear pattern visible.
```

```
...
```

```
|=====
| 50%
```

```
| Let's try some SVD now. Create the variable svd1 by assigning to it the output of a call to the
| command svd. The argument to svd should be scale(sub1[, -c(562, 563)]). This will remove the last
| columns from sub1 and scale the data. Recall that the last 2 columns contain activity and subject
| information which we won't need.
```

```
> svd1 <- svd(scale(sub1[, -c(562, 563)]))
```

```
| You are doing so well!
```

```
|=====
| 52%
```

```
| To see LEFT singular vectors of sub1, which component of svd1 would we examine?
```

```
1: x
2: u
3: d
4: v
```

```
Selection: 2
```

```
| Excellent job!
```

```
|=====
| 53%
```

```
| Call the R command dim with svd1$u as an argument.
```

```
> dim(svd1$u)
[1] 347 347
```

```
| You are quite good my friend!
```

```
|=====
| 55%
```

```
| We see that the u matrix is a 347 by 347 matrix. Each row in u corresponds to a row in the matrix
| sub1. Recall that in sub1 each row has an associated activity.
```

```
...
```

```
|=====
| 56%
```

```
| Here we're looking at the 2 left singular vectors of svd1 (the first 2 columns of svd1$u). Each
| entry of the columns belongs to a particular row with one of the 6 activities assigned to it. We
| see the activities distinguished by color. Moving from left to right, the first section of rows is
| green (standing), the second red (sitting), the third black (laying), etc. The first column of
| shows separation of the nonmoving (black, red, and green) from the walking activities. The second
| column is harder to interpret. However, the magenta cluster, which represents walking up, seems
| separate from the others.
```

```
...
```

```
|=====
| 58%
```

```
| We'll try to figure out why that is. To do that we'll have to find which of the 500+ measurements
```

```

| (represented by the columns of sub1) contributes to the variation of that component. Since we're
| interested in sub1 columns, we'll look at the RIGHT singular vectors (the columns of svd1$v), and
| in particular, the second one since the separation of the magenta cluster stood out in the second
| column of svd1$u.

...

|=====
| 59%

| Here's a plot of the second column of svd1$v. We used transparency in our plotting but nothing
| clearly stands out here. Let's use clustering to find the feature (out of the 500+) which
| contributes the most to the variation of this second column of svd1$v.

...

|=====
| 61%

| Create the variable maxCon by assigning to it the output of the R command which.max using the
| second column of svd1$v as an argument.

> maxCon <- which.max(svd1$v[,2])

| You got it right!

|=====
| 62%

| Now create a distance matrix mdist by assigning to it the output of the R command dist using 4
| columns of sub1 as the arguments. These 4 columns are 10 through 12 (10:12) and maxCon. Recall
| you'll have to concatenate these 2 column expressions when specifying them.

> mdist <- dist(c(sub1[,10:12], maxCon)
+
+
+ )
Error in dist(c(sub1[, 10:12], maxCon)) :
  (list) object cannot be coerced to type 'double'
> mdist <- dist(c(10:12, maxCon)
+ )

| Not exactly. Give it another go. Or, type info() for more options.

| Type mdist <- dist(sub1[,c(10:12,maxCon)]) at the command prompt.

> mdist <- dist(sub1[,c(10:12,maxCon)])

| You got it right!

|=====
| 64%

| Now create hclustering, the output of the R command hclust using mdist as the argument.

> hclustering <- hclust(mdist)

| You're the best!

|=====
| 66%

| Call the myplclust with 2 arguments, hclustering, and lab.col set equal to unclass(sub1$activity).

> myplclust(hclustering, lab.col = unclass(sub1$activity))

| Great job!

|=====
| 67%

```


| Now we see some real separation. Magenta (walking up) is on the far left, and the two other walking activities, the two blues, are on the far right, but in separate clusters from one another. The nonmoving activities still are jumbled together.

...

|=====

| Run the R command `names` with the argument `sub1[maxCon]` to see what measurement is associated with this maximum contributor.

```
> names(sub1[maxCon])  
[1] "fBodyAcc.meanFreq...Z"
```

| That's correct!

|=====

| So the mean body acceleration in the frequency domain in the Z direction is the main contributor to this clustering phenomenon we're seeing. Let's move on to k-means clustering to see if this technique can distinguish between the activities.

...

|=====

| Create the variable `kClust` by assigning to it the output of the R command `kmeans` with 2 arguments. The first is `sub1` with the last 2 columns removed. (Recall these don't have pertinent information for clustering analysis.) The second argument to `kmeans` is `centers` set equal to 6, the number of activities we know we have.

```
> kClust <- kmeans(sub1[,1:561], centers = 6)
```

| One more time. You can do it! Or, type `info()` for more options.

| Type `kClust <- kmeans(sub1[, -c(562, 563)], centers = 6)` the command prompt.

```
> kClust <- kmeans(sub1[, -c(562, 563)], centers = 6)
```

| You got it!

|=====

| Recall that without specifying coordinates for the cluster centroids (as we did), `kmeans` will generate starting points randomly. Here we did only 1 random start (the default). To see the output, run the R command `table` with 2 arguments. The first is `kClust$cluster` (part of the output from `kmeans`), and the second is `sub1$activity`.

```
> table  
function (..., exclude = if (useNA == "no") c(NA, NaN), useNA = c("no",  
  "ifany", "always"), dnn = list.names(...), deparse.level = 1)  
{  
  list.names <- function(...) {  
    l <- as.list(substitute(list(...)))[-1L]  
    nm <- names(l)  
    fixup <- if (is.null(nm))  
      seq_along(l)  
    else nm == ""  
    dep <- vapply(l[fixup], function(x) switch(deparse.level +  
      1, "", if (is.symbol(x)) as.character(x) else "",  
      deparse(x, nlines = 1)[1L]), "")  
    if (is.null(nm))  
      dep  
    else {  
      nm[fixup] <- dep  
      nm  
    }  
  }  
}
```

```

}
miss.use <- missing(useNA)
miss.exc <- missing(exclude)
useNA <- if (miss.use && !miss.exc && !match(NA, exclude,
  nomatch = 0L))
  "ifany"
else match.arg(useNA)
doNA <- useNA != "no"
if (!miss.use && !miss.exc && doNA && match(NA, exclude,
  nomatch = 0L))
  warning("'exclude' containing NA and 'useNA' != \"no\" are a bit contradicting")
args <- list(...)
if (!length(args))
  stop("nothing to tabulate")
if (length(args) == 1L && is.list(args[[1L]])) {
  args <- args[[1L]]
  if (length(dnn) != length(args))
    dnn <- if (!is.null(argn <- names(args)))
      argn
    else paste(dnn[1L], seq_along(args), sep = ".")
}
bin <- 0L
lens <- NULL
dims <- integer()
pd <- 1L
dn <- NULL
for (a in args) {
  if (is.null(lens))
    lens <- length(a)
  else if (length(a) != lens)
    stop("all arguments must have the same length")
  fact.a <- is.factor(a)
  if (doNA)
    aNA <- anyNA(a)
  if (!fact.a) {
    a0 <- a
    a <- factor(a, exclude = exclude)
  }
  add.na <- doNA
  if (add.na) {
    ifany <- (useNA == "ifany")
    anNAC <- anyNA(a)
    add.na <- if (!ifany || anNAC) {
      ll <- levels(a)
      if (add.ll <- !anyNA(ll)) {
        ll <- c(ll, NA)
        TRUE
      }
      else if (!ifany && !anNAC)
        FALSE
      else TRUE
    }
    else FALSE
  }
  if (add.na)
    a <- factor(a, levels = ll, exclude = NULL)
  else ll <- levels(a)
  a <- as.integer(a)
  if (fact.a && !miss.exc) {
    ll <- ll[keep <- which(match(ll, exclude, nomatch = 0L) ==
      0L)]
    a <- match(a, keep)
  }
  else if (!fact.a && add.na) {
    if (ifany && !aNA && add.ll) {
      ll <- ll[!is.na(ll)]
      is.na(a) <- match(a0, c(exclude, NA), nomatch = 0L) >
        0L
    }
    else {
      is.na(a) <- match(a0, exclude, nomatch = 0L) >

```

```

      }
    }
    n1 <- length(l1)
    dims <- c(dims, n1)
    if (prod(dims) > .Machine$integer.max)
      stop("attempt to make a table with >= 2^31 elements")
    dn <- c(dn, list(l1))
    bin <- bin + pd * (a - 1L)
    pd <- pd * n1
  }
  names(dn) <- dnn
  bin <- bin[!is.na(bin)]
  if (length(bin))
    bin <- bin + 1L
  y <- array(tabulate(bin, pd), dims, dimnames = dn)
  class(y) <- "table"
  y
}
<bytecode: 0x0000000002f6eba0>
<environment: namespace:base>

```

| You're close...I can feel it! Try it again. Or, type info() for more options.

| Type table(kClust\$cluster, sub1\$activity) the command prompt.

```
> table(kClust$cluster, sub1$activity)
```

	laying	sitting	standing	walk	walkdown	walkup
1	0	0	0	0	49	0
2	16	12	7	0	0	0
3	0	0	0	95	0	0
4	24	33	46	0	0	0
5	0	0	0	0	0	53
6	10	2	0	0	0	0

| Nice work!

```

|=====
| 75%

```

| Your exact output will depend on the state of your random number generator. We notice that when just run with 1 random start, the clusters tend to group the nonmoving activities together in one cluster. The walking activities seem to cluster individually by themselves. You could run the command to kmeans with one random start again and you'll probably get a slightly different result, but.

...

```

|=====
| 77%

```

| ... instead call kmeans with 3 arguments, the last of which will tell it to try more random starts and return the best one. The first 2 arguments should be the same as before (sub1 with the last 2 columns removed and centers set equal to 6). The third is nstart set equal to 100. Put the result in kClust again.

```
> kClust <- kmeans(sub1[, -c(562, 563)], centers = 6, nstart=100)
```

| You got it right!

```

|=====
| 78%

```

| Again, run the R command table with 2 arguments. The first is kClust\$cluster (part of the output from kmeans), and the second is sub1\$activity.

```
> table(kClust$cluster, sub1$activity)
```

	laying	sitting	standing	walk	walkdown	walkup
1	18	10	2	0	0	0
2	0	0	0	0	49	0

3	0	37	51	0	0	0
4	0	0	0	95	0	0
5	29	0	0	0	0	0
6	3	0	0	0	0	53

| Great job!

|=====

| 80%

| We see that even with 100 random starts, the passive activities tend to cluster together. One of the clusters contains only laying, but in another cluster, standing and sitting group together.

...

|=====

| 81%

| Use dim to find the dimensions of kClust's centers. Use the x\$y notation to access them.

```
> dim(kClust$centers)
[1] 6 561
```

| You're close...I can feel it! Try it again. Or, type info() for more options.

| Type dim(kClust\$centers) the command prompt.

```
> dim(kClust$centers)
[1] 6 561
```

| Excellent work!

|=====

| 83%

| So the centers are a 6 by 561 array. Sometimes it's a good idea to look at the features (columns) of these centers to see if any dominate.

...

|=====

| 84%

| Create the variable laying and assign to it the output of the call to the R command which with argument kClust\$size==29.

```
> laying <- which(kClust$size==29)
```

| All that practice is paying off!

|=====

| 86%

| Now call plot with 3 arguments. The first is kClust\$centers[laying,1:12], and the second is pch=19. The third is ylab set equal to "Laying Cluster"

```
> plot(kClust$centers[laying,1:12],pch=19, ylab="Laying Cluster")
```

| Your dedication is inspiring!

|=====

| 88%

| We see the first 3 columns dominate this cluster center. Run names with the first 3 columns of s as the argument to remind yourself of what these columns contain.

```
> names(sub1[,1:3])
[1] "tBodyAcc.mean...X" "tBodyAcc.mean...Y" "tBodyAcc.mean...Z"
```

| Keep up the great work!

|=====|
| 89%

| So the 3 directions of mean body acceleration seem to have the biggest effect on laying.

...

|=====|
| 91%

| Create the variable walkdown and assign to it the output of the call to the R command which with
| the argument kClust\$size==49.

```
> walkdown <- which(kClust$size==49)
```

| You nailed it! Good job!

|=====|
| 92%

| Now call plot with 3 arguments. The first is kClust\$centers[walkdown,1:12], and the second is pch
| set to 19. The third is ylab set equal to "walkdown cluster"

```
> plot(kClust$centers[walkdown,1:12],pch=19, ylab="walkdown cluster")
```

| You got it!

|=====|
| 94%

| We see an interesting pattern here. From left to right, looking at the 12 acceleration measurements
| in groups of 3, the points decrease in value. The X direction dominates, followed by Y then Z. This
| might tell us something more about the walking down activity.

...

|=====|
| 95%

| We'll wrap up here and hope this example convinced you that real world analysis can be frustrating
| sometimes and not always obvious. You might have to try several techniques of exploratory data
| analysis before you hit one that pays off and leads you to the questions that will be the most
| promising to explore.

...

==|=====|
| 97%

| We saw here that the sensor measurements were pretty good at discriminating between the 3 walking
| activities, but the passive activities were harder to distinguish from one another. These might
| require more analysis or an entirely different set of sensory measurements.

...

===|=====|
| 98%

| Congratulations! We hope you enjoyed the 6 activities and 500+ features of this lesson.

...

====|=====|
| 100%

| Would you like to receive credit for completing this course on Coursera.org?

1: No
2: Yes

Codes

```
showXY.R
```

```
par(mfrow=c(1, 2), mar = c(5, 4, 1, 1))
```

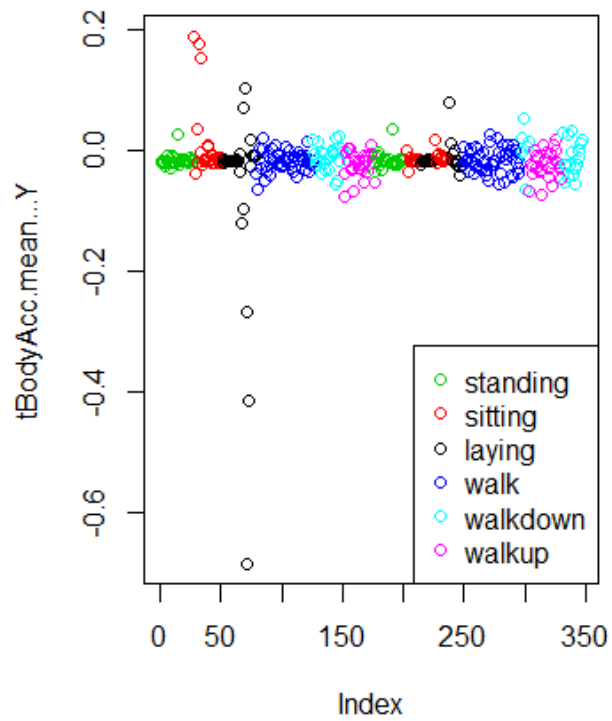

```
plot(sub1[, 1], col = sub1$activity, ylab = names(sub1)[1])
```



```
plot(sub1[, 2], col = sub1$activity, ylab = names(sub1)[2])
```



```
legend("bottomright", legend=unique(sub1$activity), col=unique(sub1$activity), pch = 1)
```

```
par(mfrow=c(1,1))
```

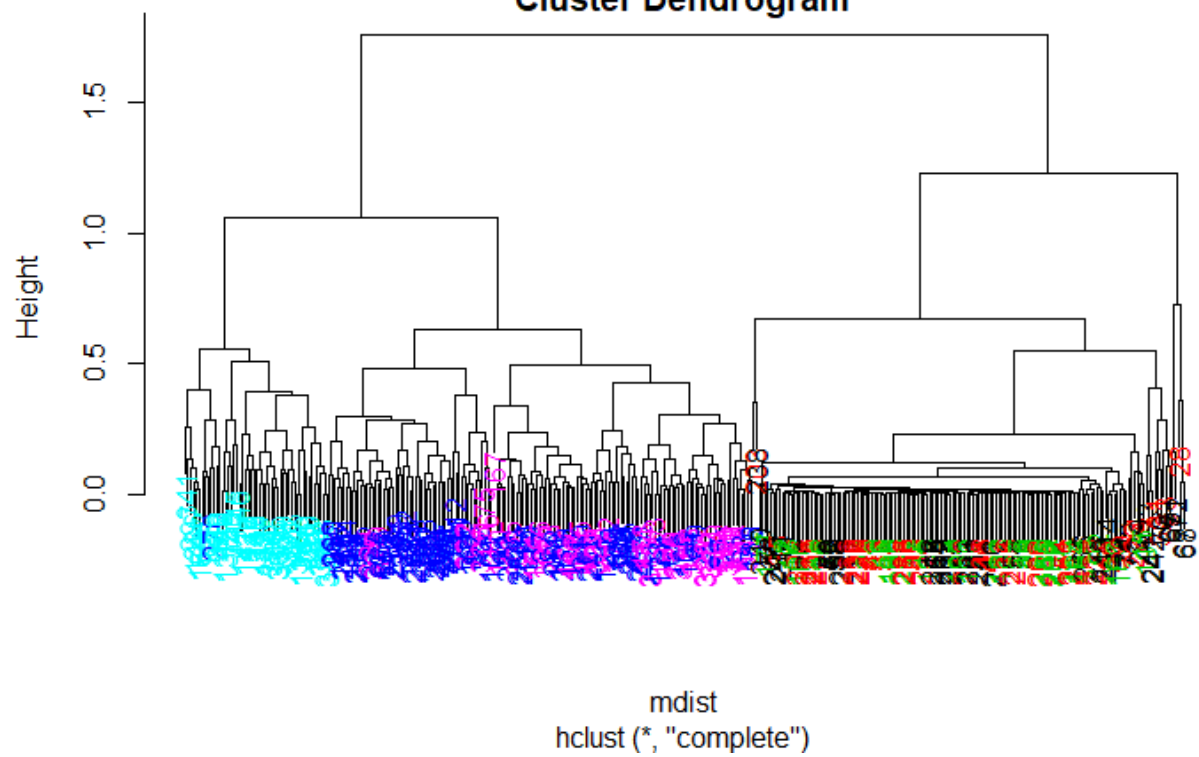
Plots








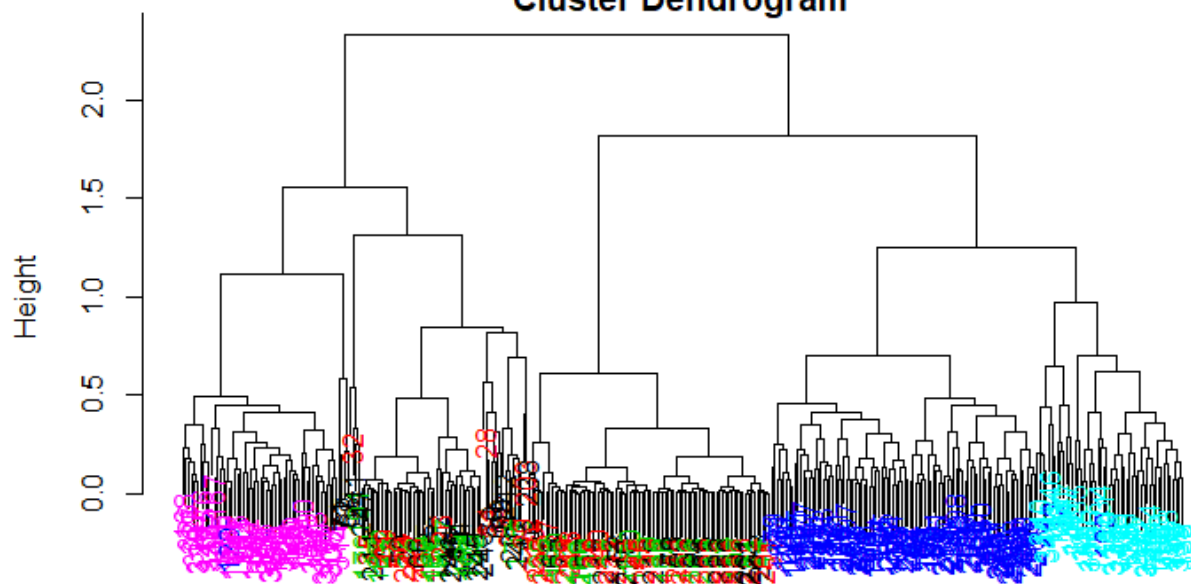


Cluster Dendrogram





Cluster Dendrogram



mdist
hclust (*, "complete")

